

Vivek Convolution Networking Assignment

Report: Exploring the Relationship Between Training Sample Size and Choice of Network for Cats vs. Dogs Classification

Objective

The goal of this project was to investigate the impact of training sample size on the performance of convolutional neural networks (CNNs) trained from scratch versus using pretrained networks. Specifically, we aimed to:

1. Apply CNNs to image data and explore overfitting reduction techniques.
2. Compare the performance of models trained from scratch and pretrained models across different training sample sizes (1500, 2000, and 2500).
3. Optimize model architectures and techniques to achieve the best performance.

Methodology

I trained two types of models on a subset of the Cats vs. Dogs dataset:

1. Scratch Model: A CNN model built from scratch with multiple convolutional layers, pooling layers, and dropout regularization to reduce overfitting.
2. Pretrained Model: A transfer learning approach using the ResNet50 architecture pretrained on the ImageNet dataset, with fine-tuning applied on the final layers.

Both models were trained on three different sample sizes: 1500, 2000, and 2500 training samples, with 500 validation and 500 test samples for each scenario. Techniques such as data augmentation, dropout, and regularization were applied to mitigate overfitting.

Results

The following table summarizes the performance of both models across different sample sizes:

Sample Size	Model Type	Training Accuracy	Validation Accuracy	Training Loss	Validation Loss
1500	Scratch	0.6007	0.562	1.8124	1.8398
1500	Pretrained	0.9473	0.966	0.1224	0.0985
2000	Scratch	0.6375	0.552	1.7352	2.1645
2000	Pretrained	0.9415	0.958	0.1472	0.1302
2500	Scratch	0.6264	0.616	1.3995	1.3633
2500	Pretrained	0.9448	0.966	0.1357	0.0932

Analysis and Key Findings

1. Performance of Scratch Models:

- The scratch models exhibited lower training and validation accuracy compared to the pretrained models across all sample sizes. The best validation accuracy achieved by the scratch model was 0.616 for the 2500 sample size.
- The models trained from scratch were prone to overfitting, as indicated by a significant gap between training accuracy and validation accuracy, particularly with the 2000 and 2500 sample sizes.
- Regularization techniques (such as dropout and data augmentation) were effective at mitigating overfitting, but their overall impact was limited compared to the performance of pretrained models.

2. Performance of Pretrained Models:

- Pretrained models consistently outperformed the scratch models, achieving validation accuracies of 0.958 to 0.966 across all sample sizes, with minimal overfitting. The validation losses were also significantly lower than those of the scratch models.
- Pretrained models demonstrated that even with smaller datasets (1500 samples), transfer learning provided excellent generalization, confirming that pretrained networks are highly effective for small to moderately sized datasets.

3. Impact of Sample Size:

- Increasing the training sample size improved the performance of both scratch and pretrained models, but the improvement was marginal for pretrained models since they already achieved high accuracy even with smaller sample sizes.

- Scratch models benefitted slightly more from larger sample sizes, but still underperformed compared to pretrained models. The training process for scratch models appeared to stabilize around the 2500 sample size but did not close the performance gap.

4. Overfitting and Regularization:

- Overfitting was more pronounced in scratch models, which had lower validation accuracy and higher validation loss compared to training loss. This suggests that scratch models required larger datasets and additional regularization techniques to improve generalization.

- Pretrained models, benefiting from their pretrained weights, were less susceptible to overfitting and required fewer overfitting prevention techniques.

Conclusion

In conclusion, this project highlights a clear relationship between training sample size and the choice of network:

- **Pretrained Models:** Transfer learning with pretrained models (like ResNet50) is highly effective, especially when working with small to moderate sample sizes. Pretrained models are able to generalize well with less data and require less fine-tuning compared to scratch models.

- **Scratch Models:** While scratch models can achieve reasonable performance, they require much larger datasets to avoid overfitting and reach the same level of generalization as pretrained models. For smaller datasets, scratch models are prone to overfitting and underperformance, even with regularization techniques applied.

Based on these findings, pretrained models are recommended when working with smaller datasets due to their superior performance and generalization ability. Scratch models may be suitable for larger datasets or when a custom architecture is required, but they demand more careful tuning and regularization.

Loading Libraries

```
In [1]: from google.colab import drive
import zipfile
import os
import shutil
import random
import matplotlib.pyplot as plt
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense, Dropout, BatchNormalization
from tensorflow.keras.regularizers import l2
from tensorflow.keras.applications import InceptionV3
from tensorflow.keras.layers import GlobalAveragePooling2D
from tensorflow.keras.models import Model
```

Loading Dataset

```
In [2]: # Mount Google Drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

```
In [3]: # Define paths
base_dir = '/content/drive/MyDrive'
zip_file_path = os.path.join(base_dir, 'cats_vs_dogs_small_dataset.zip')
extracted_dir_path = os.path.join(base_dir, 'cats_vs_dogs_small_dataset')
```

```
In [4]: # Unzip dataset
with zipfile.ZipFile(zip_file_path, 'r') as zip_ref:
    zip_ref.extractall(extracted_dir_path)
```

```
In [4]: # Dataset directories for 'cat' and 'dog' folders
cat_folder_path = os.path.join(extracted_dir_path, 'cats_vs_dogs_small_dataset/cat')
dog_folder_path = os.path.join(extracted_dir_path, 'cats_vs_dogs_small_dataset/dog')
```

Splitting The Dataset

```
In [5]: # Function to split and organize dataset
def split_data_folders(base_dir, cat_images, dog_images, train_samples, validation_samples):
    train_dir = os.path.join(base_dir, 'train')
    validation_dir = os.path.join(base_dir, 'validation')
    test_dir = os.path.join(base_dir, 'test')

    for d in [train_dir, validation_dir, test_dir]:
        shutil.rmtree(d, ignore_errors=True)
        os.makedirs(os.path.join(d, 'cat'), exist_ok=True)
        os.makedirs(os.path.join(d, 'dog'), exist_ok=True)
```

```

random.shuffle(cat_images)
random.shuffle(dog_images)

def copy_images(src_dir, dst_dir, file_list):
    for file in file_list:
        src_path = os.path.join(src_dir, file)
        dst_path = os.path.join(dst_dir, file)
        shutil.copyfile(src_path, dst_path)

copy_images(cat_folder_path, os.path.join(train_dir, 'cat'), cat_images[:train_samples])
copy_images(dog_folder_path, os.path.join(train_dir, 'dog'), dog_images[:train_samples])

copy_images(cat_folder_path, os.path.join(validation_dir, 'cat'),
             cat_images[train_samples // 2:train_samples // 2 + validation_samples])
copy_images(dog_folder_path, os.path.join(validation_dir, 'dog'),
             dog_images[train_samples // 2:train_samples // 2 + validation_samples])

copy_images(cat_folder_path, os.path.join(test_dir, 'cat'),
             cat_images[train_samples // 2 + validation_samples // 2:
                        train_samples // 2 + validation_samples // 2 + test_samples])
copy_images(dog_folder_path, os.path.join(test_dir, 'dog'),
             dog_images[train_samples // 2 + validation_samples // 2:
                        train_samples // 2 + validation_samples // 2 + test_samples])

return train_dir, validation_dir, test_dir

```

```

In [6]: # Image augmentations and data generators
def create_data_generators(train_dir, validation_dir, test_dir, image_size, batch_size):
    train_datagen = ImageDataGenerator(
        rescale=1./255,
        rotation_range=40,
        width_shift_range=0.2,
        height_shift_range=0.2,
        shear_range=0.2,
        zoom_range=0.2,
        horizontal_flip=True,
        fill_mode='nearest'
    )
    validation_datagen = ImageDataGenerator(rescale=1./255)
    test_datagen = ImageDataGenerator(rescale=1./255)

    train_generator = train_datagen.flow_from_directory(
        train_dir,
        target_size=image_size,
        batch_size=batch_size,
        class_mode='binary'
    )
    validation_generator = validation_datagen.flow_from_directory(
        validation_dir,
        target_size=image_size,
        batch_size=batch_size,
        class_mode='binary'
    )
    test_generator = test_datagen.flow_from_directory(
        test_dir,
        target_size=image_size,
        batch_size=batch_size,
        class_mode='binary'
    )

```

```
)
return train_generator, validation_generator, test_generator
```

Building Optimized Scratch Model

```
In [7]: # Improved Scratch Model
def build_optimized_scratch_model(image_size):
    model = Sequential([
        Conv2D(64, (3, 3), activation='relu', input_shape=(image_size[0], image_size[1], 3),
        BatchNormalization(),
        MaxPooling2D((2, 2)),
        Conv2D(128, (3, 3), activation='relu', kernel_regularizer=l2(0.001)),
        BatchNormalization(),
        MaxPooling2D((2, 2)),
        Conv2D(256, (3, 3), activation='relu', kernel_regularizer=l2(0.001)),
        BatchNormalization(),
        MaxPooling2D((2, 2)),
        Conv2D(512, (3, 3), activation='relu', kernel_regularizer=l2(0.001)),
        BatchNormalization(),
        MaxPooling2D((2, 2)),
        Flatten(),
        Dense(512, activation='relu', kernel_regularizer=l2(0.001)),
        Dropout(0.5), # Prevent overfitting
        Dense(1, activation='sigmoid')
    ])
    model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
    return model
```

Inception Model

```
In [8]: # Pretrained InceptionV3 Model
def build_inception_pretrained_model(image_size):
    base_model = InceptionV3(weights='imagenet', include_top=False, input_shape=(image_size[0], image_size[1], 3))
    for layer in base_model.layers:
        layer.trainable = False # Freeze the base model layers
    x = GlobalAveragePooling2D()(base_model.output)
    x = Dense(512, activation='relu')(x)
    Dropout(0.5)
    output = Dense(1, activation='sigmoid')(x)
    model = Model(inputs=base_model.input, outputs=output)
    model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
    return model
```

```
In [9]: # Train the model and plot results
def train_and_evaluate_model(model, train_generator, validation_generator, epochs):
    history = model.fit(
        train_generator,
        epochs=epochs,
        validation_data=validation_generator
    )
    return history
```

```
In [10]: # Plot performance metrics
def plot_training_metrics(history):
```

```

plt.plot(history.history['accuracy'], label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.title('Training and Validation Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()
plt.show()

plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.title('Training and Validation Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()
plt.show()

```

```

In [23]: # Plot accuracy comparison for scratch vs pretrained model for different sample sizes
def plot_accuracy_comparison(results, sample_size):
    # Access the history attributes correctly
    scratch_acc = results[sample_size]['scratch']['history'].history['accuracy']
    scratch_val_acc = results[sample_size]['scratch']['history'].history['val_accuracy']
    pretrained_acc = results[sample_size]['pretrained']['history'].history['accuracy']
    pretrained_val_acc = results[sample_size]['pretrained']['history'].history['val_ac

    plt.plot(scratch_acc, label='Scratch Model Training Accuracy')
    plt.plot(scratch_val_acc, label='Scratch Model Validation Accuracy')
    plt.plot(pretrained_acc, label='Pretrained Model Training Accuracy')
    plt.plot(pretrained_val_acc, label='Pretrained Model Validation Accuracy')
    plt.title(f'Training vs Validation Accuracy for {sample_size} Samples')
    plt.xlabel('Epochs')
    plt.ylabel('Accuracy')
    plt.legend()
    plt.show()

```

```

In [24]: # Plot Loss comparison for scratch vs pretrained model for different sample sizes
def plot_loss_comparison(results, sample_size):
    # Access the history attributes correctly
    scratch_loss = results[sample_size]['scratch']['history'].history['loss']
    scratch_val_loss = results[sample_size]['scratch']['history'].history['val_loss']
    pretrained_loss = results[sample_size]['pretrained']['history'].history['loss']
    pretrained_val_loss = results[sample_size]['pretrained']['history'].history['val_l

    plt.plot(scratch_loss, label='Scratch Model Training Loss')
    plt.plot(scratch_val_loss, label='Scratch Model Validation Loss')
    plt.plot(pretrained_loss, label='Pretrained Model Training Loss')
    plt.plot(pretrained_val_loss, label='Pretrained Model Validation Loss')
    plt.title(f'Training vs Validation Loss for {sample_size} Samples')
    plt.xlabel('Epochs')
    plt.ylabel('Loss')
    plt.legend()
    plt.show()

```

```

In [15]: # Setting up parameters
image_size = (150, 150)
batch_size = 32
validation_samples = 500
test_samples = 500
cat_images = os.listdir(cat_folder_path)
dog_images = os.listdir(dog_folder_path)

```

```
In [16]: # Train and validate models for both 1500 and 2000 sample sizes
sample_sizes = [1500, 2000, 2500]
results = {}
```

```
In [17]: scratch_model = build_optimized_scratch_model(image_size)
scratch_model.summary()
```

```
/usr/local/lib/python3.10/dist-packages/keras/src/layers/convolutional/base_conv.py:107: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.
  super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 148, 148, 64)	
batch_normalization (BatchNormalization)	(None, 148, 148, 64)	
max_pooling2d (MaxPooling2D)	(None, 74, 74, 64)	
conv2d_1 (Conv2D)	(None, 72, 72, 128)	
batch_normalization_1 (BatchNormalization)	(None, 72, 72, 128)	
max_pooling2d_1 (MaxPooling2D)	(None, 36, 36, 128)	
conv2d_2 (Conv2D)	(None, 34, 34, 256)	
batch_normalization_2 (BatchNormalization)	(None, 34, 34, 256)	
max_pooling2d_2 (MaxPooling2D)	(None, 17, 17, 256)	
conv2d_3 (Conv2D)	(None, 15, 15, 512)	
batch_normalization_3 (BatchNormalization)	(None, 15, 15, 512)	
max_pooling2d_3 (MaxPooling2D)	(None, 7, 7, 512)	
flatten (Flatten)	(None, 25088)	
dense (Dense)	(None, 512)	
dropout (Dropout)	(None, 512)	
dense_1 (Dense)	(None, 1)	

Total params: 14,400,897 (54.94 MB)

Trainable params: 14,398,977 (54.93 MB)


```
In [18]: pretrained_model = build_inception_pretrained_model(image_size)
pretrained_model.summary()
```

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/inception_v3/inception_v3_weights_tf_dim_ordering_tf_kernels_notop.h5

87910968/87910968 ————— 5s 0us/step

Model: "functional_1"



Layer (type)	Output Shape	Param #	Connected to
input_layer_1 (InputLayer)	(None, 150, 150, 3)	0	-
conv2d_4 (Conv2D)	(None, 74, 74, 32)	864	input_layer_1
batch_normalization_4 (BatchNormalization)	(None, 74, 74, 32)	96	conv2d_4
activation (Activation)	(None, 74, 74, 32)	0	batch_normalization_4
conv2d_5 (Conv2D)	(None, 72, 72, 32)	9,216	activation
batch_normalization_5 (BatchNormalization)	(None, 72, 72, 32)	96	conv2d_5
activation_1 (Activation)	(None, 72, 72, 32)	0	batch_normalization_5
conv2d_6 (Conv2D)	(None, 72, 72, 64)	18,432	activation_1
batch_normalization_6 (BatchNormalization)	(None, 72, 72, 64)	192	conv2d_6
activation_2 (Activation)	(None, 72, 72, 64)	0	batch_normalization_6
max_pooling2d_4 (MaxPooling2D)	(None, 35, 35, 64)	0	activation_2
conv2d_7 (Conv2D)	(None, 35, 35, 80)	5,120	max_pooling2d_4
batch_normalization_7 (BatchNormalization)	(None, 35, 35, 80)	240	conv2d_7
activation_3 (Activation)	(None, 35, 35, 80)	0	batch_normalization_7
conv2d_8 (Conv2D)	(None, 33, 33, 192)	138,240	activation_3
batch_normalization_8 (BatchNormalization)	(None, 33, 33, 192)	576	conv2d_8
activation_4 (Activation)	(None, 33, 33, 192)	0	batch_normalization_8
max_pooling2d_5 (MaxPooling2D)	(None, 16, 16, 192)	0	activation_4
conv2d_12 (Conv2D)	(None, 16, 16, 64)	12,288	max_pooling2d_5
batch_normalization_12 (BatchNormalization)	(None, 16, 16, 64)	192	conv2d_12
activation_8 (Activation)	(None, 16, 16, 64)	0	batch_normalization_12
conv2d_10 (Conv2D)	(None, 16, 16, 48)	9,216	activation_8
conv2d_13 (Conv2D)	(None, 16, 16, 96)	55,296	conv2d_10

batch_normalization_10 (BatchNormalization)	(None, 16, 16, 48)	144	conv2d_16
batch_normalization_13 (BatchNormalization)	(None, 16, 16, 96)	288	conv2d_17
activation_6 (Activation)	(None, 16, 16, 48)	0	batch_norm
activation_9 (Activation)	(None, 16, 16, 96)	0	batch_norm
average_pooling2d (AveragePooling2D)	(None, 16, 16, 192)	0	max_pooli
conv2d_9 (Conv2D)	(None, 16, 16, 64)	12,288	max_pooli
conv2d_11 (Conv2D)	(None, 16, 16, 64)	76,800	activation
conv2d_14 (Conv2D)	(None, 16, 16, 96)	82,944	activation
conv2d_15 (Conv2D)	(None, 16, 16, 32)	6,144	average_p
batch_normalization_9 (BatchNormalization)	(None, 16, 16, 64)	192	conv2d_9[
batch_normalization_11 (BatchNormalization)	(None, 16, 16, 64)	192	conv2d_11
batch_normalization_14 (BatchNormalization)	(None, 16, 16, 96)	288	conv2d_14
batch_normalization_15 (BatchNormalization)	(None, 16, 16, 32)	96	conv2d_15
activation_5 (Activation)	(None, 16, 16, 64)	0	batch_norm
activation_7 (Activation)	(None, 16, 16, 64)	0	batch_norm
activation_10 (Activation)	(None, 16, 16, 96)	0	batch_norm
activation_11 (Activation)	(None, 16, 16, 32)	0	batch_norm
mixed0 (Concatenate)	(None, 16, 16, 256)	0	activation activation activation activation
conv2d_19 (Conv2D)	(None, 16, 16, 64)	16,384	mixed0[0]
batch_normalization_19 (BatchNormalization)	(None, 16, 16, 64)	192	conv2d_19
activation_15 (Activation)	(None, 16, 16, 64)	0	batch_norm

conv2d_17 (Conv2D)	(None , 16 , 16 , 48)	12,288	mixed0[0]
conv2d_20 (Conv2D)	(None , 16 , 16 , 96)	55,296	activation_13
batch_normalization_17 (BatchNormalization)	(None , 16 , 16 , 48)	144	conv2d_17
batch_normalization_20 (BatchNormalization)	(None , 16 , 16 , 96)	288	conv2d_20
activation_13 (Activation)	(None , 16 , 16 , 48)	0	batch_normalization_17
activation_16 (Activation)	(None , 16 , 16 , 96)	0	batch_normalization_20
average_pooling2d_1 (AveragePooling2D)	(None , 16 , 16 , 256)	0	mixed0[0]
conv2d_16 (Conv2D)	(None , 16 , 16 , 64)	16,384	mixed0[0]
conv2d_18 (Conv2D)	(None , 16 , 16 , 64)	76,800	activation_16
conv2d_21 (Conv2D)	(None , 16 , 16 , 96)	82,944	activation_18
conv2d_22 (Conv2D)	(None , 16 , 16 , 64)	16,384	average_pooling2d_1
batch_normalization_16 (BatchNormalization)	(None , 16 , 16 , 64)	192	conv2d_16
batch_normalization_18 (BatchNormalization)	(None , 16 , 16 , 64)	192	conv2d_18
batch_normalization_21 (BatchNormalization)	(None , 16 , 16 , 96)	288	conv2d_21
batch_normalization_22 (BatchNormalization)	(None , 16 , 16 , 64)	192	conv2d_22
activation_12 (Activation)	(None , 16 , 16 , 64)	0	batch_normalization_16
activation_14 (Activation)	(None , 16 , 16 , 64)	0	batch_normalization_18
activation_17 (Activation)	(None , 16 , 16 , 96)	0	batch_normalization_21
activation_18 (Activation)	(None , 16 , 16 , 64)	0	batch_normalization_22
mixed1 (Concatenate)	(None , 16 , 16 , 288)	0	activation_12 activation_14 activation_17 activation_18
conv2d_26 (Conv2D)	(None , 16 , 16 , 64)	18,432	mixed1[0]

batch_normalization_26 (BatchNormalization)	(None, 16, 16, 64)	192	conv2d_26
activation_22 (Activation)	(None, 16, 16, 64)	0	batch_norm
conv2d_24 (Conv2D)	(None, 16, 16, 48)	13,824	mixed1[0]
conv2d_27 (Conv2D)	(None, 16, 16, 96)	55,296	activation_
batch_normalization_24 (BatchNormalization)	(None, 16, 16, 48)	144	conv2d_24
batch_normalization_27 (BatchNormalization)	(None, 16, 16, 96)	288	conv2d_27
activation_20 (Activation)	(None, 16, 16, 48)	0	batch_norm
activation_23 (Activation)	(None, 16, 16, 96)	0	batch_norm
average_pooling2d_2 (AveragePooling2D)	(None, 16, 16, 288)	0	mixed1[0]
conv2d_23 (Conv2D)	(None, 16, 16, 64)	18,432	mixed1[0]
conv2d_25 (Conv2D)	(None, 16, 16, 64)	76,800	activation_
conv2d_28 (Conv2D)	(None, 16, 16, 96)	82,944	activation_
conv2d_29 (Conv2D)	(None, 16, 16, 64)	18,432	average_p
batch_normalization_23 (BatchNormalization)	(None, 16, 16, 64)	192	conv2d_23
batch_normalization_25 (BatchNormalization)	(None, 16, 16, 64)	192	conv2d_25
batch_normalization_28 (BatchNormalization)	(None, 16, 16, 96)	288	conv2d_28
batch_normalization_29 (BatchNormalization)	(None, 16, 16, 64)	192	conv2d_29
activation_19 (Activation)	(None, 16, 16, 64)	0	batch_norm
activation_21 (Activation)	(None, 16, 16, 64)	0	batch_norm
activation_24 (Activation)	(None, 16, 16, 96)	0	batch_norm
activation_25 (Activation)	(None, 16, 16, 64)	0	batch_norm

mixed2 (Concatenate)	(None, 16, 16, 288)	0	activation activation activation activation
conv2d_31 (Conv2D)	(None, 16, 16, 64)	18,432	mixed2[0]
batch_normalization_31 (BatchNormalization)	(None, 16, 16, 64)	192	conv2d_31
activation_27 (Activation)	(None, 16, 16, 64)	0	batch_norm
conv2d_32 (Conv2D)	(None, 16, 16, 96)	55,296	activation
batch_normalization_32 (BatchNormalization)	(None, 16, 16, 96)	288	conv2d_32
activation_28 (Activation)	(None, 16, 16, 96)	0	batch_norm
conv2d_30 (Conv2D)	(None, 7, 7, 384)	995,328	mixed2[0]
conv2d_33 (Conv2D)	(None, 7, 7, 96)	82,944	activation
batch_normalization_30 (BatchNormalization)	(None, 7, 7, 384)	1,152	conv2d_30
batch_normalization_33 (BatchNormalization)	(None, 7, 7, 96)	288	conv2d_33
activation_26 (Activation)	(None, 7, 7, 384)	0	batch_norm
activation_29 (Activation)	(None, 7, 7, 96)	0	batch_norm
max_pooling2d_6 (MaxPooling2D)	(None, 7, 7, 288)	0	mixed2[0]
mixed3 (Concatenate)	(None, 7, 7, 768)	0	activation activation max_pooli
conv2d_38 (Conv2D)	(None, 7, 7, 128)	98,304	mixed3[0]
batch_normalization_38 (BatchNormalization)	(None, 7, 7, 128)	384	conv2d_38
activation_34 (Activation)	(None, 7, 7, 128)	0	batch_norm
conv2d_39 (Conv2D)	(None, 7, 7, 128)	114,688	activation
batch_normalization_39 (BatchNormalization)	(None, 7, 7, 128)	384	conv2d_39

activation_35 (Activation)	(None, 7, 7, 128)	0	batch_norm
conv2d_35 (Conv2D)	(None, 7, 7, 128)	98,304	mixed3[0]
conv2d_40 (Conv2D)	(None, 7, 7, 128)	114,688	activation
batch_normalization_35 (BatchNormalization)	(None, 7, 7, 128)	384	conv2d_35
batch_normalization_40 (BatchNormalization)	(None, 7, 7, 128)	384	conv2d_40
activation_31 (Activation)	(None, 7, 7, 128)	0	batch_norm
activation_36 (Activation)	(None, 7, 7, 128)	0	batch_norm
conv2d_36 (Conv2D)	(None, 7, 7, 128)	114,688	activation
conv2d_41 (Conv2D)	(None, 7, 7, 128)	114,688	activation
batch_normalization_36 (BatchNormalization)	(None, 7, 7, 128)	384	conv2d_36
batch_normalization_41 (BatchNormalization)	(None, 7, 7, 128)	384	conv2d_41
activation_32 (Activation)	(None, 7, 7, 128)	0	batch_norm
activation_37 (Activation)	(None, 7, 7, 128)	0	batch_norm
average_pooling2d_3 (AveragePooling2D)	(None, 7, 7, 768)	0	mixed3[0]
conv2d_34 (Conv2D)	(None, 7, 7, 192)	147,456	mixed3[0]
conv2d_37 (Conv2D)	(None, 7, 7, 192)	172,032	activation
conv2d_42 (Conv2D)	(None, 7, 7, 192)	172,032	activation
conv2d_43 (Conv2D)	(None, 7, 7, 192)	147,456	average_p
batch_normalization_34 (BatchNormalization)	(None, 7, 7, 192)	576	conv2d_34
batch_normalization_37 (BatchNormalization)	(None, 7, 7, 192)	576	conv2d_37
batch_normalization_42 (BatchNormalization)	(None, 7, 7, 192)	576	conv2d_42
batch_normalization_43	(None, 7, 7, 192)	576	conv2d_43

(BatchNormalization)			
activation_30 (Activation)	(None, 7, 7, 192)	0	batch_norm
activation_33 (Activation)	(None, 7, 7, 192)	0	batch_norm
activation_38 (Activation)	(None, 7, 7, 192)	0	batch_norm
activation_39 (Activation)	(None, 7, 7, 192)	0	batch_norm
mixed4 (Concatenate)	(None, 7, 7, 768)	0	activation_39 activation_38 activation_33 activation_30
conv2d_48 (Conv2D)	(None, 7, 7, 160)	122,880	mixed4[0]
batch_normalization_48 (BatchNormalization)	(None, 7, 7, 160)	480	conv2d_48
activation_44 (Activation)	(None, 7, 7, 160)	0	batch_norm
conv2d_49 (Conv2D)	(None, 7, 7, 160)	179,200	activation_44
batch_normalization_49 (BatchNormalization)	(None, 7, 7, 160)	480	conv2d_49
activation_45 (Activation)	(None, 7, 7, 160)	0	batch_norm
conv2d_45 (Conv2D)	(None, 7, 7, 160)	122,880	mixed4[0]
conv2d_50 (Conv2D)	(None, 7, 7, 160)	179,200	activation_45
batch_normalization_45 (BatchNormalization)	(None, 7, 7, 160)	480	conv2d_50
batch_normalization_50 (BatchNormalization)	(None, 7, 7, 160)	480	conv2d_45
activation_41 (Activation)	(None, 7, 7, 160)	0	batch_norm
activation_46 (Activation)	(None, 7, 7, 160)	0	batch_norm
conv2d_46 (Conv2D)	(None, 7, 7, 160)	179,200	activation_46
conv2d_51 (Conv2D)	(None, 7, 7, 160)	179,200	activation_41
batch_normalization_46 (BatchNormalization)	(None, 7, 7, 160)	480	conv2d_51

batch_normalization_51 (BatchNormalization)	(None, 7, 7, 160)	480	conv2d_51
activation_42 (Activation)	(None, 7, 7, 160)	0	batch_norm
activation_47 (Activation)	(None, 7, 7, 160)	0	batch_norm
average_pooling2d_4 (AveragePooling2D)	(None, 7, 7, 768)	0	mixed4[0]
conv2d_44 (Conv2D)	(None, 7, 7, 192)	147,456	mixed4[0]
conv2d_47 (Conv2D)	(None, 7, 7, 192)	215,040	activation
conv2d_52 (Conv2D)	(None, 7, 7, 192)	215,040	activation
conv2d_53 (Conv2D)	(None, 7, 7, 192)	147,456	average_p
batch_normalization_44 (BatchNormalization)	(None, 7, 7, 192)	576	conv2d_44
batch_normalization_47 (BatchNormalization)	(None, 7, 7, 192)	576	conv2d_47
batch_normalization_52 (BatchNormalization)	(None, 7, 7, 192)	576	conv2d_52
batch_normalization_53 (BatchNormalization)	(None, 7, 7, 192)	576	conv2d_53
activation_40 (Activation)	(None, 7, 7, 192)	0	batch_norm
activation_43 (Activation)	(None, 7, 7, 192)	0	batch_norm
activation_48 (Activation)	(None, 7, 7, 192)	0	batch_norm
activation_49 (Activation)	(None, 7, 7, 192)	0	batch_norm
mixed5 (Concatenate)	(None, 7, 7, 768)	0	activation activation activation activation
conv2d_58 (Conv2D)	(None, 7, 7, 160)	122,880	mixed5[0]
batch_normalization_58 (BatchNormalization)	(None, 7, 7, 160)	480	conv2d_58
activation_54 (Activation)	(None, 7, 7, 160)	0	batch_norm

conv2d_59 (Conv2D)	(None, 7, 7, 160)	179,200	activation_59
batch_normalization_59 (BatchNormalization)	(None, 7, 7, 160)	480	conv2d_59
activation_55 (Activation)	(None, 7, 7, 160)	0	batch_normalization_59
conv2d_55 (Conv2D)	(None, 7, 7, 160)	122,880	mixed5[0]
conv2d_60 (Conv2D)	(None, 7, 7, 160)	179,200	activation_55
batch_normalization_55 (BatchNormalization)	(None, 7, 7, 160)	480	conv2d_55
batch_normalization_60 (BatchNormalization)	(None, 7, 7, 160)	480	conv2d_60
activation_51 (Activation)	(None, 7, 7, 160)	0	batch_normalization_60
activation_56 (Activation)	(None, 7, 7, 160)	0	batch_normalization_51
conv2d_56 (Conv2D)	(None, 7, 7, 160)	179,200	activation_56
conv2d_61 (Conv2D)	(None, 7, 7, 160)	179,200	activation_56
batch_normalization_56 (BatchNormalization)	(None, 7, 7, 160)	480	conv2d_56
batch_normalization_61 (BatchNormalization)	(None, 7, 7, 160)	480	conv2d_61
activation_52 (Activation)	(None, 7, 7, 160)	0	batch_normalization_61
activation_57 (Activation)	(None, 7, 7, 160)	0	batch_normalization_52
average_pooling2d_5 (AveragePooling2D)	(None, 7, 7, 768)	0	mixed5[0]
conv2d_54 (Conv2D)	(None, 7, 7, 192)	147,456	mixed5[0]
conv2d_57 (Conv2D)	(None, 7, 7, 192)	215,040	activation_54
conv2d_62 (Conv2D)	(None, 7, 7, 192)	215,040	activation_57
conv2d_63 (Conv2D)	(None, 7, 7, 192)	147,456	average_pooling2d_5
batch_normalization_54 (BatchNormalization)	(None, 7, 7, 192)	576	conv2d_54
batch_normalization_57 (BatchNormalization)	(None, 7, 7, 192)	576	conv2d_57

batch_normalization_62 (BatchNormalization)	(None, 7, 7, 192)	576	conv2d_62
batch_normalization_63 (BatchNormalization)	(None, 7, 7, 192)	576	conv2d_63
activation_50 (Activation)	(None, 7, 7, 192)	0	batch_norm
activation_53 (Activation)	(None, 7, 7, 192)	0	batch_norm
activation_58 (Activation)	(None, 7, 7, 192)	0	batch_norm
activation_59 (Activation)	(None, 7, 7, 192)	0	batch_norm
mixed6 (Concatenate)	(None, 7, 7, 768)	0	activation activation activation activation
conv2d_68 (Conv2D)	(None, 7, 7, 192)	147,456	mixed6[0]
batch_normalization_68 (BatchNormalization)	(None, 7, 7, 192)	576	conv2d_68
activation_64 (Activation)	(None, 7, 7, 192)	0	batch_norm
conv2d_69 (Conv2D)	(None, 7, 7, 192)	258,048	activation
batch_normalization_69 (BatchNormalization)	(None, 7, 7, 192)	576	conv2d_69
activation_65 (Activation)	(None, 7, 7, 192)	0	batch_norm
conv2d_65 (Conv2D)	(None, 7, 7, 192)	147,456	mixed6[0]
conv2d_70 (Conv2D)	(None, 7, 7, 192)	258,048	activation
batch_normalization_65 (BatchNormalization)	(None, 7, 7, 192)	576	conv2d_65
batch_normalization_70 (BatchNormalization)	(None, 7, 7, 192)	576	conv2d_70
activation_61 (Activation)	(None, 7, 7, 192)	0	batch_norm
activation_66 (Activation)	(None, 7, 7, 192)	0	batch_norm
conv2d_66 (Conv2D)	(None, 7, 7, 192)	258,048	activation

conv2d_71 (Conv2D)	(None , 7, 7, 192)	258,048	activation_61
batch_normalization_66 (BatchNormalization)	(None , 7, 7, 192)	576	conv2d_66
batch_normalization_71 (BatchNormalization)	(None , 7, 7, 192)	576	conv2d_71
activation_62 (Activation)	(None , 7, 7, 192)	0	batch_normalization_66
activation_67 (Activation)	(None , 7, 7, 192)	0	batch_normalization_71
average_pooling2d_6 (AveragePooling2D)	(None , 7, 7, 768)	0	mixed6[0]
conv2d_64 (Conv2D)	(None , 7, 7, 192)	147,456	mixed6[0]
conv2d_67 (Conv2D)	(None , 7, 7, 192)	258,048	activation_62
conv2d_72 (Conv2D)	(None , 7, 7, 192)	258,048	activation_67
conv2d_73 (Conv2D)	(None , 7, 7, 192)	147,456	average_pooling2d_6
batch_normalization_64 (BatchNormalization)	(None , 7, 7, 192)	576	conv2d_64
batch_normalization_67 (BatchNormalization)	(None , 7, 7, 192)	576	conv2d_67
batch_normalization_72 (BatchNormalization)	(None , 7, 7, 192)	576	conv2d_72
batch_normalization_73 (BatchNormalization)	(None , 7, 7, 192)	576	conv2d_73
activation_60 (Activation)	(None , 7, 7, 192)	0	batch_normalization_64
activation_63 (Activation)	(None , 7, 7, 192)	0	batch_normalization_67
activation_68 (Activation)	(None , 7, 7, 192)	0	batch_normalization_72
activation_69 (Activation)	(None , 7, 7, 192)	0	batch_normalization_73
mixed7 (Concatenate)	(None , 7, 7, 768)	0	activation_60 activation_63 activation_68 activation_69
conv2d_76 (Conv2D)	(None , 7, 7, 192)	147,456	mixed7[0]

batch_normalization_76 (BatchNormalization)	(None, 7, 7, 192)	576	conv2d_76
activation_72 (Activation)	(None, 7, 7, 192)	0	batch_norm
conv2d_77 (Conv2D)	(None, 7, 7, 192)	258,048	activation
batch_normalization_77 (BatchNormalization)	(None, 7, 7, 192)	576	conv2d_77
activation_73 (Activation)	(None, 7, 7, 192)	0	batch_norm
conv2d_74 (Conv2D)	(None, 7, 7, 192)	147,456	mixed7[0]
conv2d_78 (Conv2D)	(None, 7, 7, 192)	258,048	activation
batch_normalization_74 (BatchNormalization)	(None, 7, 7, 192)	576	conv2d_74
batch_normalization_78 (BatchNormalization)	(None, 7, 7, 192)	576	conv2d_78
activation_70 (Activation)	(None, 7, 7, 192)	0	batch_norm
activation_74 (Activation)	(None, 7, 7, 192)	0	batch_norm
conv2d_75 (Conv2D)	(None, 3, 3, 320)	552,960	activation
conv2d_79 (Conv2D)	(None, 3, 3, 192)	331,776	activation
batch_normalization_75 (BatchNormalization)	(None, 3, 3, 320)	960	conv2d_75
batch_normalization_79 (BatchNormalization)	(None, 3, 3, 192)	576	conv2d_79
activation_71 (Activation)	(None, 3, 3, 320)	0	batch_norm
activation_75 (Activation)	(None, 3, 3, 192)	0	batch_norm
max_pooling2d_7 (MaxPooling2D)	(None, 3, 3, 768)	0	mixed7[0]
mixed8 (Concatenate)	(None, 3, 3, 1280)	0	activation activation max_pooli
conv2d_84 (Conv2D)	(None, 3, 3, 448)	573,440	mixed8[0]
batch_normalization_84 (BatchNormalization)	(None, 3, 3, 448)	1,344	conv2d_84

activation_80 (Activation)	(None, 3, 3, 448)	0	batch_norm
conv2d_81 (Conv2D)	(None, 3, 3, 384)	491,520	mixed8[0]
conv2d_85 (Conv2D)	(None, 3, 3, 384)	1,548,288	activation
batch_normalization_81 (BatchNormalization)	(None, 3, 3, 384)	1,152	conv2d_81
batch_normalization_85 (BatchNormalization)	(None, 3, 3, 384)	1,152	conv2d_85
activation_77 (Activation)	(None, 3, 3, 384)	0	batch_norm
activation_81 (Activation)	(None, 3, 3, 384)	0	batch_norm
conv2d_82 (Conv2D)	(None, 3, 3, 384)	442,368	activation
conv2d_83 (Conv2D)	(None, 3, 3, 384)	442,368	activation
conv2d_86 (Conv2D)	(None, 3, 3, 384)	442,368	activation
conv2d_87 (Conv2D)	(None, 3, 3, 384)	442,368	activation
average_pooling2d_7 (AveragePooling2D)	(None, 3, 3, 1280)	0	mixed8[0]
conv2d_80 (Conv2D)	(None, 3, 3, 320)	409,600	mixed8[0]
batch_normalization_82 (BatchNormalization)	(None, 3, 3, 384)	1,152	conv2d_82
batch_normalization_83 (BatchNormalization)	(None, 3, 3, 384)	1,152	conv2d_83
batch_normalization_86 (BatchNormalization)	(None, 3, 3, 384)	1,152	conv2d_86
batch_normalization_87 (BatchNormalization)	(None, 3, 3, 384)	1,152	conv2d_87
conv2d_88 (Conv2D)	(None, 3, 3, 192)	245,760	average_p
batch_normalization_80 (BatchNormalization)	(None, 3, 3, 320)	960	conv2d_80
activation_78 (Activation)	(None, 3, 3, 384)	0	batch_norm
activation_79 (Activation)	(None, 3, 3, 384)	0	batch_norm
activation_82	(None, 3, 3, 384)	0	batch_norm

(Activation)			
activation_83 (Activation)	(None, 3, 3, 384)	0	batch_norm
batch_normalization_88 (BatchNormalization)	(None, 3, 3, 192)	576	conv2d_88
activation_76 (Activation)	(None, 3, 3, 320)	0	batch_norm
mixed9_0 (Concatenate)	(None, 3, 3, 768)	0	activation_76 activation_76
concatenate (Concatenate)	(None, 3, 3, 768)	0	activation_76 activation_76
activation_84 (Activation)	(None, 3, 3, 192)	0	batch_norm
mixed9 (Concatenate)	(None, 3, 3, 2048)	0	activation_84 mixed9_0[0] concatenate activation_84
conv2d_93 (Conv2D)	(None, 3, 3, 448)	917,504	mixed9[0]
batch_normalization_93 (BatchNormalization)	(None, 3, 3, 448)	1,344	conv2d_93
activation_89 (Activation)	(None, 3, 3, 448)	0	batch_norm
conv2d_90 (Conv2D)	(None, 3, 3, 384)	786,432	mixed9[0]
conv2d_94 (Conv2D)	(None, 3, 3, 384)	1,548,288	activation_89
batch_normalization_90 (BatchNormalization)	(None, 3, 3, 384)	1,152	conv2d_90
batch_normalization_94 (BatchNormalization)	(None, 3, 3, 384)	1,152	conv2d_94
activation_86 (Activation)	(None, 3, 3, 384)	0	batch_norm
activation_90 (Activation)	(None, 3, 3, 384)	0	batch_norm
conv2d_91 (Conv2D)	(None, 3, 3, 384)	442,368	activation_86
conv2d_92 (Conv2D)	(None, 3, 3, 384)	442,368	activation_90
conv2d_95 (Conv2D)	(None, 3, 3, 384)	442,368	activation_90
conv2d_96 (Conv2D)	(None, 3, 3, 384)	442,368	activation_90

average_pooling2d_8 (AveragePooling2D)	(None, 3, 3, 2048)	0	mixed9[0]
conv2d_89 (Conv2D)	(None, 3, 3, 320)	655,360	mixed9[0]
batch_normalization_91 (BatchNormalization)	(None, 3, 3, 384)	1,152	conv2d_90
batch_normalization_92 (BatchNormalization)	(None, 3, 3, 384)	1,152	conv2d_90
batch_normalization_95 (BatchNormalization)	(None, 3, 3, 384)	1,152	conv2d_90
batch_normalization_96 (BatchNormalization)	(None, 3, 3, 384)	1,152	conv2d_90
conv2d_97 (Conv2D)	(None, 3, 3, 192)	393,216	average_p
batch_normalization_89 (BatchNormalization)	(None, 3, 3, 320)	960	conv2d_89
activation_87 (Activation)	(None, 3, 3, 384)	0	batch_nor
activation_88 (Activation)	(None, 3, 3, 384)	0	batch_nor
activation_91 (Activation)	(None, 3, 3, 384)	0	batch_nor
activation_92 (Activation)	(None, 3, 3, 384)	0	batch_nor
batch_normalization_97 (BatchNormalization)	(None, 3, 3, 192)	576	conv2d_97
activation_85 (Activation)	(None, 3, 3, 320)	0	batch_nor
mixed9_1 (Concatenate)	(None, 3, 3, 768)	0	activation activation
concatenate_1 (Concatenate)	(None, 3, 3, 768)	0	activation activation
activation_93 (Activation)	(None, 3, 3, 192)	0	batch_nor
mixed10 (Concatenate)	(None, 3, 3, 2048)	0	activation mixed9_1[concatena activation
global_average_pooling2d (GlobalAveragePooling2D)	(None, 2048)	0	mixed10[0]

dense_2 (Dense)	(None, 512)	1,049,088	global_av
dense_3 (Dense)	(None, 1)	513	dense_2[6

Total params: 22,852,385 (87.17 MB)

Trainable params: 1,049,601 (4.00 MB)

Model Training

```
In [20]: for sample_size in sample_sizes:
# Split dataset for current sample size
train_dir, validation_dir, test_dir = split_data_folders(extracted_dir_path, cat_i
train_generator, validation_generator, test_generator = create_data_generators(tr

# Train scratch model
scratch_model = build_optimized_scratch_model(image_size)
history_scratch = train_and_evaluate_model(scratch_model, train_generator, validat
test_loss_scratch, test_accuracy_scratch = scratch_model.evaluate(test_generator)

# Train pretrained InceptionV3 model
pretrained_model = build_inception_pretrained_model(image_size)
history_pretrained = train_and_evaluate_model(pretrained_model, train_generator, v
test_loss_pretrained, test_accuracy_pretrained = pretrained_model.evaluate(test_ge

# Store results for later comparison
results[sample_size] = {
    'scratch': {'model': scratch_model, 'history': history_scratch, 'test_loss': t
    'pretrained': {'model': pretrained_model, 'history': history_pretrained, 'test

}

# Plot training and validation performance
print(f"\nResults for CNN Model {sample_size} samples:")
plot_training_metrics(history_scratch)
print(f"\nResults for Pre Trained Model {sample_size} samples:")
plot_training_metrics(history_pretrained)
```


Found 1500 images belonging to 2 classes.


Found 500 images belonging to 2 classes.


Found 500 images belonging to 2 classes.


```
/usr/local/lib/python3.10/dist-packages/keras/src/layers/convolutional/base_conv.py:1
07: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When u
sing Sequential models, prefer using an `Input(shape)` object as the first layer in t
he model instead.
```


```
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```


Epoch 1/10
47/47  32s 482ms/step - accuracy: 0.5472 - loss: 7.4070 - val_accuracy: 0.5020 - val_loss: 6.7758


Epoch 2/10
47/47  20s 362ms/step - accuracy: 0.5454 - loss: 5.8942 - val_accuracy: 0.5080 - val_loss: 4.1834


Epoch 3/10
47/47  19s 351ms/step - accuracy: 0.5697 - loss: 3.6160 - val_accuracy: 0.5260 - val_loss: 3.0504


Epoch 4/10
47/47  20s 368ms/step - accuracy: 0.5357 - loss: 2.6757 - val_accuracy: 0.5040 - val_loss: 3.3322



Epoch 5/10
47/47  19s 355ms/step - accuracy: 0.5310 - loss: 2.4724 - val_accuracy: 0.6080 - val_loss: 2.3365


Epoch 6/10
47/47  21s 364ms/step - accuracy: 0.5584 - loss: 2.3234 - val_accuracy: 0.5440 - val_loss: 2.2301


Epoch 7/10
47/47  20s 350ms/step - accuracy: 0.5928 - loss: 2.1722 - val_accuracy: 0.5080 - val_loss: 3.0899


Epoch 8/10
47/47  20s 359ms/step - accuracy: 0.6407 - loss: 2.0324 - val_accuracy: 0.5740 - val_loss: 2.0331


Epoch 9/10
47/47  22s 422ms/step - accuracy: 0.6032 - loss: 1.9424 - val_accuracy: 0.5840 - val_loss: 1.8505


Epoch 10/10
47/47  41s 420ms/step - accuracy: 0.6155 - loss: 1.8153 - val_accuracy: 0.5620 - val_loss: 1.8398
16/16  3s 163ms/step - accuracy: 0.5343 - loss: 1.8164


Epoch 1/10
47/47  61s 943ms/step - accuracy: 0.7164 - loss: 2.2800 - val_accuracy: 0.9680 - val_loss: 0.0916


Epoch 2/10
47/47  21s 404ms/step - accuracy: 0.9205 - loss: 0.1786 - val_accuracy: 0.9640 - val_loss: 0.1082


Epoch 3/10
47/47  19s 358ms/step - accuracy: 0.9159 - loss: 0.2050 - val_accuracy: 0.9620 - val_loss: 0.0905


Epoch 4/10
47/47  23s 415ms/step - accuracy: 0.9263 - loss: 0.1524 - val_accuracy: 0.9520 - val_loss: 0.0908


Epoch 5/10
47/47  19s 354ms/step - accuracy: 0.9374 - loss: 0.1591 - val_accuracy: 0.9620 - val_loss: 0.0844

Epoch 6/10
47/47  23s 417ms/step - accuracy: 0.9214 - loss: 0.1613 - val_accuracy: 0.9680 - val_loss: 0.0863

Epoch 7/10
47/47  19s 358ms/step - accuracy: 0.9408 - loss: 0.1312 - val_accuracy: 0.9640 - val_loss: 0.1005

Epoch 8/10
47/47  19s 360ms/step - accuracy: 0.9451 - loss: 0.1290 - val_accuracy: 0.9560 - val_loss: 0.1121


Epoch 9/10
47/47  20s 343ms/step - accuracy: 0.9544 - loss: 0.1147 - val_accuracy: 0.9580 - val_loss: 0.0883


Epoch 10/10
47/47  18s 339ms/step - accuracy: 0.9530 - loss: 0.1092 - val_accuracy: 0.9530 - val_loss: 0.1092

uracy: 0.9660 - val_loss: 0.0985


16/16 ————— 2s 149ms/step - accuracy: 0.9601 - loss: 0.0988


Results for CNN Model 1500 samples:





Results for Pre Trained Model 1500 samples:






Found 2000 images belonging to 2 classes.


Found 500 images belonging to 2 classes.

Found 500 images belonging to 2 classes.


Epoch 1/10

63/63  **42s** 519ms/step - accuracy: 0.5371 - loss: 9.3586 - val_accuracy: 0.5000 - val_loss: 5.8278


Epoch 2/10

63/63  **29s** 363ms/step - accuracy: 0.5154 - loss: 4.8557 - val_accuracy: 0.5000 - val_loss: 4.5087


Epoch 3/10

63/63  **39s** 351ms/step - accuracy: 0.5600 - loss: 2.9523 - val_accuracy: 0.5220 - val_loss: 2.5862


Epoch 4/10

63/63  **42s** 361ms/step - accuracy: 0.6022 - loss: 2.5199 - val_accuracy: 0.5040 - val_loss: 2.8582


Epoch 5/10

63/63  **25s** 356ms/step - accuracy: 0.6036 - loss: 2.3387 - val_accuracy: 0.4800 - val_loss: 2.8316


Epoch 6/10

63/63  **26s** 365ms/step - accuracy: 0.6289 - loss: 2.1685 - val_accuracy: 0.5580 - val_loss: 2.1522


Epoch 7/10

63/63  **25s** 364ms/step - accuracy: 0.6198 - loss: 2.0263 - val_accuracy: 0.5720 - val_loss: 2.0322


Epoch 8/10

63/63  **41s** 365ms/step - accuracy: 0.6357 - loss: 1.8809 - val_accuracy: 0.5360 - val_loss: 2.2022

Epoch 9/10


63/63  **24s** 342ms/step - accuracy: 0.6400 - loss: 1.7775 - val_accuracy: 0.5440 - val_loss: 2.4304

Epoch 10/10


63/63  **24s** 337ms/step - accuracy: 0.6430 - loss: 1.7204 - val_accuracy: 0.5520 - val_loss: 2.1645

16/16  **2s** 138ms/step - accuracy: 0.6269 - loss: 1.9761


Epoch 1/10

63/63  **48s** 506ms/step - accuracy: 0.8193 - loss: 0.7808 - val_accuracy: 0.9460 - val_loss: 0.1630


Epoch 2/10

63/63  **24s** 350ms/step - accuracy: 0.9277 - loss: 0.1984 - val_accuracy: 0.9600 - val_loss: 0.1108


Epoch 3/10

63/63  **23s** 332ms/step - accuracy: 0.9316 - loss: 0.1667 - val_accuracy: 0.9440 - val_loss: 0.1672


Epoch 4/10

63/63  **25s** 351ms/step - accuracy: 0.9430 - loss: 0.1343 - val_accuracy: 0.9280 - val_loss: 0.1747


Epoch 5/10

63/63  **24s** 350ms/step - accuracy: 0.9056 - loss: 0.2149 - val_accuracy: 0.9380 - val_loss: 0.1887


Epoch 6/10

63/63  **43s** 386ms/step - accuracy: 0.9339 - loss: 0.1473 - val_accuracy: 0.9440 - val_loss: 0.1384

Epoch 7/10

63/63  **26s** 373ms/step - accuracy: 0.9485 - loss: 0.1252 - val_accuracy: 0.9540 - val_loss: 0.1103

Epoch 8/10

63/63  **23s** 327ms/step - accuracy: 0.9397 - loss: 0.1436 - val_accuracy: 0.9660 - val_loss: 0.0969

Epoch 9/10

63/63  **24s** 338ms/step - accuracy: 0.9462 - loss: 0.1346 - val_accuracy: 0.9462 - val_loss: 0.1346

uracy: 0.9620 - val_loss: 0.1094


Epoch 10/10


63/63 ————— **41s** 334ms/step - accuracy: 0.9426 - loss: 0.1443 - val_acc

uracy: 0.9580 - val_loss: 0.1302


16/16 ————— **3s** 190ms/step - accuracy: 0.9621 - loss: 0.1031


Results for CNN Model 2000 samples:





Results for Pre Trained Model 2000 samples:






Found 2500 images belonging to 2 classes.


Found 500 images belonging to 2 classes.

Found 500 images belonging to 2 classes.


Epoch 1/10

79/79  **46s** 475ms/step - accuracy: 0.5313 - loss: 6.8356 - val_accuracy: 0.4680 - val_loss: 3.3398


Epoch 2/10

79/79  **31s** 367ms/step - accuracy: 0.5419 - loss: 3.7844 - val_accuracy: 0.4940 - val_loss: 2.6578


Epoch 3/10

79/79  **31s** 366ms/step - accuracy: 0.5785 - loss: 2.4775 - val_accuracy: 0.5000 - val_loss: 2.4924


Epoch 4/10

79/79  **31s** 367ms/step - accuracy: 0.6107 - loss: 2.2325 - val_accuracy: 0.5040 - val_loss: 2.9432


Epoch 5/10

79/79  **31s** 360ms/step - accuracy: 0.6394 - loss: 2.0145 - val_accuracy: 0.5640 - val_loss: 1.9163


Epoch 6/10

79/79  **40s** 359ms/step - accuracy: 0.6264 - loss: 1.8425 - val_accuracy: 0.5720 - val_loss: 2.0305


Epoch 7/10

79/79  **33s** 373ms/step - accuracy: 0.5675 - loss: 2.0198 - val_accuracy: 0.5700 - val_loss: 1.9591


Epoch 8/10

79/79  **41s** 379ms/step - accuracy: 0.6313 - loss: 1.7120 - val_accuracy: 0.5860 - val_loss: 1.6105

Epoch 9/10


79/79  **31s** 361ms/step - accuracy: 0.6317 - loss: 1.5143 - val_accuracy: 0.6260 - val_loss: 1.4046

Epoch 10/10


79/79  **41s** 363ms/step - accuracy: 0.6450 - loss: 1.3951 - val_accuracy: 0.6160 - val_loss: 1.3633

16/16  **2s** 146ms/step - accuracy: 0.6096 - loss: 1.3484


Epoch 1/10

79/79  **54s** 507ms/step - accuracy: 0.7597 - loss: 1.9802 - val_accuracy: 0.9580 - val_loss: 0.0997


Epoch 2/10

79/79  **30s** 350ms/step - accuracy: 0.9270 - loss: 0.1758 - val_accuracy: 0.9600 - val_loss: 0.1051


Epoch 3/10

79/79  **30s** 352ms/step - accuracy: 0.9276 - loss: 0.1675 - val_accuracy: 0.9560 - val_loss: 0.1132


Epoch 4/10

79/79  **41s** 345ms/step - accuracy: 0.9450 - loss: 0.1422 - val_accuracy: 0.9620 - val_loss: 0.0888


Epoch 5/10

79/79  **29s** 349ms/step - accuracy: 0.9338 - loss: 0.1544 - val_accuracy: 0.9620 - val_loss: 0.1071


Epoch 6/10

79/79  **30s** 348ms/step - accuracy: 0.9348 - loss: 0.1516 - val_accuracy: 0.9580 - val_loss: 0.1040

Epoch 7/10

79/79  **41s** 344ms/step - accuracy: 0.9419 - loss: 0.1516 - val_accuracy: 0.9560 - val_loss: 0.1044

Epoch 8/10

79/79  **41s** 343ms/step - accuracy: 0.9485 - loss: 0.1343 - val_accuracy: 0.9600 - val_loss: 0.1222

Epoch 9/10

79/79  **29s** 339ms/step - accuracy: 0.9363 - loss: 0.1484 - val_accuracy: 0.9363 - val_loss: 0.1484

uracy: 0.9640 - val_loss: 0.0942

Epoch 10/10


79/79 ————— **41s** 334ms/step - accuracy: 0.9410 - loss: 0.1451 - val_acc

uracy: 0.9660 - val_loss: 0.0932


16/16 ————— **3s** 159ms/step - accuracy: 0.9840 - loss: 0.0584


Results for CNN Model 2500 samples:





Results for Pre Trained Model 2500 samples:








Model Comparison

```
In [25]: # Plot accuracy and loss comparison for both models for all sample sizes
for sample_size in sample_sizes:
    print(f"\nResults for {sample_size} samples:")
    plot_accuracy_comparison(results, sample_size)
    plot_loss_comparison(results, sample_size)
```



Results for 1500 samples:



Results for 2000 samples:



Results for 2500 samples:



```
In [30]: import pandas as pd
def display_results_table(results, sample_sizes):
```



```

# Prepare lists to store the data
data = {
    'Sample Size': [],
    'Model Type': [],
    'Training Accuracy': [],
    'Validation Accuracy': [],
    'Training Loss': [],
    'Validation Loss': []
}

for sample_size in sample_sizes:
    # Scratch Model Results
    scratch_acc = results[sample_size]['scratch']['history'].history['accuracy'][-1]
    scratch_val_acc = results[sample_size]['scratch']['history'].history['val_accuracy'][-1]
    scratch_loss = results[sample_size]['scratch']['history'].history['loss'][-1]
    scratch_val_loss = results[sample_size]['scratch']['history'].history['val_loss'][-1]

    # Pretrained Model Results
    pretrained_acc = results[sample_size]['pretrained']['history'].history['accuracy'][-1]
    pretrained_val_acc = results[sample_size]['pretrained']['history'].history['val_accuracy'][-1]
    pretrained_loss = results[sample_size]['pretrained']['history'].history['loss'][-1]
    pretrained_val_loss = results[sample_size]['pretrained']['history'].history['val_loss'][-1]

    # Append data for scratch model
    data['Sample Size'].append(sample_size)
    data['Model Type'].append('Scratch')
    data['Training Accuracy'].append(scratch_acc)
    data['Validation Accuracy'].append(scratch_val_acc)
    data['Training Loss'].append(scratch_loss)
    data['Validation Loss'].append(scratch_val_loss)

    # Append data for pretrained model
    data['Sample Size'].append(sample_size)
    data['Model Type'].append('Pretrained')
    data['Training Accuracy'].append(pretrained_acc)
    data['Validation Accuracy'].append(pretrained_val_acc)
    data['Training Loss'].append(pretrained_loss)
    data['Validation Loss'].append(pretrained_val_loss)

# Create a DataFrame
df = pd.DataFrame(data)

# Print the DataFrame
return df

```

```




In [31]: df = display_results_table(results, sample_sizes)
df

```


Out[31]:

	Sample Size	Model Type	Training Accuracy	Validation Accuracy	Training Loss	Validation Loss
0	1500	Scratch	0.600667	0.562	1.812383	1.839815
1	1500	Pretrained	0.947333	0.966	0.122357	0.098470
2	2000	Scratch	0.637500	0.552	1.735171	2.164541
3	2000	Pretrained	0.941500	0.958	0.147226	0.130177
4	2500	Scratch	0.626400	0.616	1.399527	1.363309
5	2500	Pretrained	0.944800	0.966	0.135690	0.093156





Distributions





Categorical distributions




2-d distributions

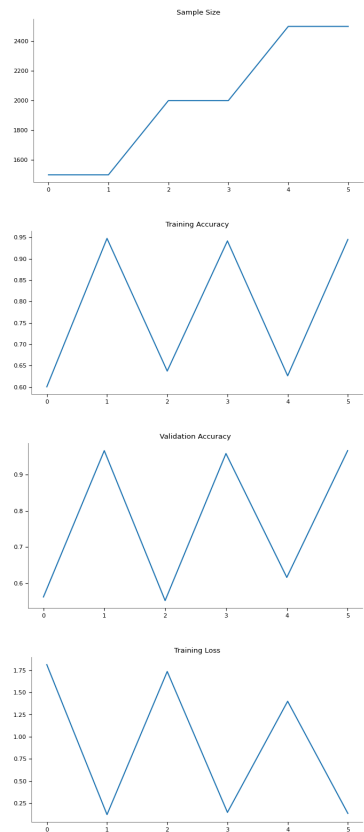


Time series






Values



Faceted distributions

<string>:5: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `y` variable to `hue` and set `legend=False` for the same effect.



<string>:5: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `y` variable to `hue` and set `legend=False` for the same effect.



<string>:5: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `y` variable to `hue` and set `legend=False` for the same effect.



<string>:5: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `y` variable to `hue` and set `legend=False` for the same effect.

